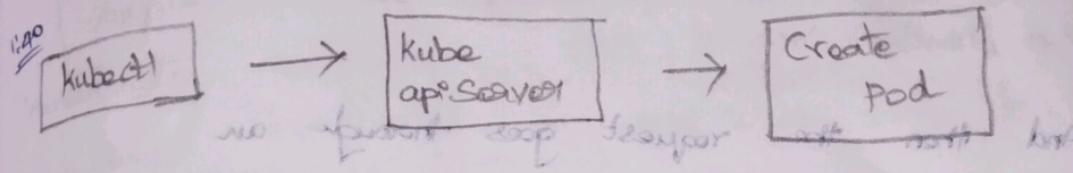
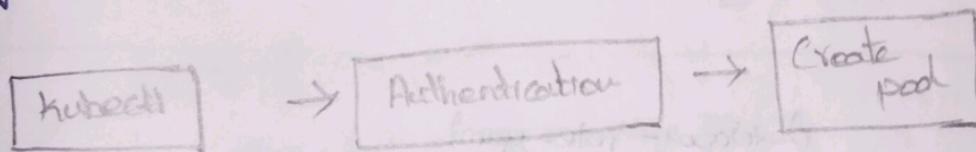


83 (2025 updates) Admission Controllers



⇒ we've been running commands from our command line using the "kubectl" utility to perform various kinds of operations on our k8s cluster. It's just like

⇒ Every time when we send a request, say, to create a pod, the request goes to the API server & then the pod is created & the information is finally persisted in the Etcd database.



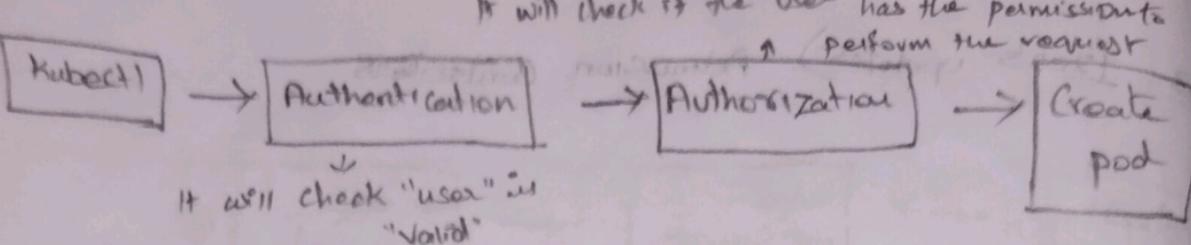
⇒ When the requests hit the API server, we've learned that it goes through an authentication process & this is usually done through certificates

⇒ If the request was sent through kubectl, we know the Kube Config file has the Certificates configured & the authentication process is responsible for identifying the user who sent the request & making sure the user is valid. [The API server has a kubeconfig file which has a certificate which contains process to check whether the user is valid]

> cat .kube/config

apiVersion: v1
 - clusters:
 - <cluster>
 - <cluster>
 - <cluster>
 - <cluster>
 - <cluster>

server: https://10.10.10.15:6443



And then the request goes through an

"authorization process" this is when we check by the user has permission to perform that operation

this is achieved through rbac-role-based access control, simpler a base on roles and verbs

\Rightarrow So if a user is assigned a particular role of a "developer", the user is allowed to $*$ list $*$ update, $*$ get, $*$ create, $*$ delete pods

developer-role.yaml

apiVersion: rbac.authorization.k8s.io/v1
kind: Role

metadata:

name: developer

rules:

- apiGroups: [""]

resources: ["pods"]

verbs: ["list", "create", "get", "update", "delete"]

\Rightarrow And so, if the request that came in matched any of these conditions, in this case it does a request is to create a pod, it is allowed to go through, otherwise it is rejected. So this is "authorization" with "role-based access control"

Authorization - RBAC 133

With RBAC, you could place in different kinds of restriction. Such as "to allow" or "to deny" those with a particular role.

* To create a list (or)

* To delete different kinds of objects like pods, deployments (or) services

* We could even restrict access to specific resource names, such as a developer can only work on pods named "blue" (or) "orange" (or) restrict access within a namespace alone

developer-role.yaml

apiVersion: rbac.authorization.k8s.io/v1

kind: Role

metadata:

name: developer

rules:

- apiGroups: [" "] : resources

resources: ["pods"]

verbs: ["create"]

resourceNames: ["blue", "orange"]

* can list pods, deployments / services / ...

* can create pods / deployments / services / ...

* can delete pods / deployments / services / ...

* can create pods named blue or orange

* can create pods within a namespace

→ Most of the rules that you can create with "role-based access control" is @ the k8s API level
↳ what user is allowed, access to what kind of API operations & it doesn't go beyond that

↳ what if we want to do more than just define what kind of access a user has to an object

↳ what if we want to do more than just define what kind of access a user has to an object

apiVersion: v1

kind: Pod

metadata:
name: web-pod

spec:

containers:

- name: ubuntu
image: ubuntu-latest
command: ["sleep", "3600"]

securityContext:

runAsUser: 0

capabilities:

["egress", "add": ["MAC=ADMIN"]]

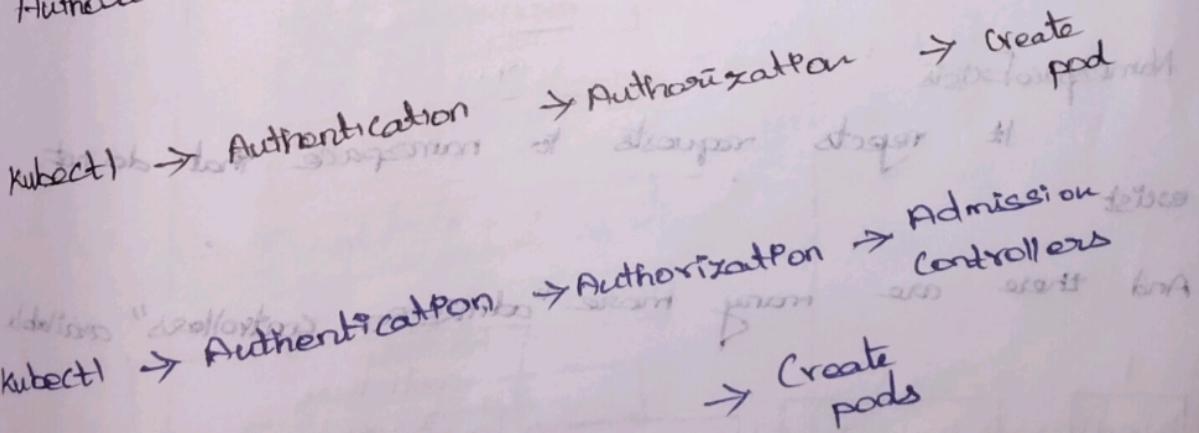
for e.g:-

When a pod creation request comes in & you'd like to review the configuration file & look at the image name & say you do not want to allow images from a public docker

hub registry, only allows images from a specific hub.
 Internal registry, only permit images from certain registered.
 Only permit images from certain registry.
 the latest tag for any images. For e.g..

- If you say that container is running as the root user, then you do not want to allow that request (or) runAsUser: 0 [Do not permit runAs root user] grants only certain capabilities only certain capabilities (or) to enforce things.
- Allow certain capabilities only certain capabilities (or) to enforce things.
- that the metadata always contains labels add: ["MAC-ADMIN"] {Pods always has labels}
- * Only permit images from certain registry.
- * Do not permit runAs root user
- * Only permit certain capabilities.
- * Pod always has labels.

These are the some of the things that you cannot achieve with the existing RBAC & this is whole "Admission Controllers" comes in.



⇒ Admission controllers help us implement better security measures to enforce how a cluster is used.

⇒ Apart from simply validating configuration, admission controllers can do a lot more. Suggest change the request itself (or) perform additional operations before the pod gets created.

⇒ There are number of "admission controllers" that come pre-built with k8s such as **Always**.

pull images that ensures that everytime a pod is created the images are always pulled.

DefaultStorageClass :-

It observes creation of PVs & automatically adds a default storage class to them. If one is not specified,

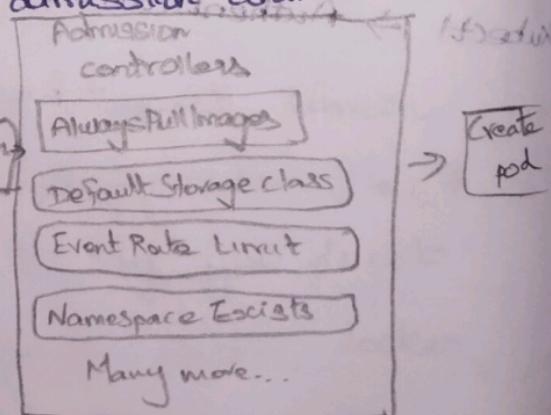
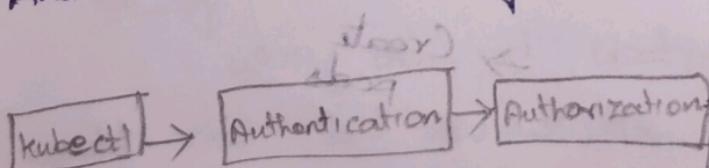
EventRateLimit :-

It can help set a limit on the requests with the API server can handle at a time to prevent the API server from flooding with requests.

NamespaceExists

It rejects requests to namespace that do not exist.

And there are many more "admission controllers" available.



E.g.: The NamespaceExists admission controller

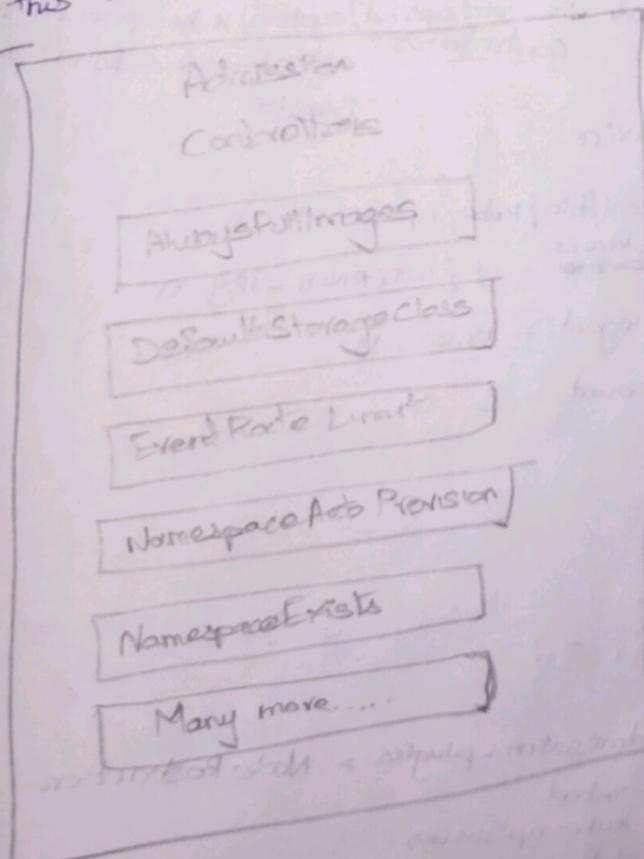
→ Say we want to create a pod in a namespace called "Blue" that doesn't exist

Kubectl run nginx --image nginx --namespace blue.
Error from server (NotFound): namespace "blue" not found.

What's happening here is, that my request gets authenticated, then authorized & it then goes through the admission controllers, "NamespaceExists" admission controller handles the request & checks if the "blue" namespace is available. If it is not, the request is rejected.

The "namespace exists" is a built-in admission controller that is enabled by default.

There is another admission controller that is not enabled by default, & that is called as Namespace Auto Provision Admission Controller. This will automatically create the namespace, if it doesn't exist.



View Enabled Admission Controllers

→ To see list of admission controllers enabled by default
run `kube-apiserver -h` command

`kube-apiserver -h | grep enable-admission-plugins`

The command will list admission controllers that are enabled by default, the ones that are highlighted in "green"

NOTE: If you're running this in a Kubernetes based setup then you must run this command ~~within the~~ Kube API Server control plane pod using the kubectl exec command first like this

`sudo kubectl exec kube-apiserver-control-plane -n kube-system`

`-- kube-apiserver -h | grep enable-admission-plugins`

⇒ sudo vim /etc/systemd/system/kube-apiserver.service
Controllers add "admissionController" under "ExecStart"

Enable Admission Kube-apiserver.service

ExecStart = /usr/local/bin/kube-apiserver //

-- advertise-~~address~~ = \${INTERNAL-IP} //

-- allow-Privileged = true //

-- apiserver-count = 3 //

This approach is for Kube API Server Service

-- enable-admission-plugins = NodeRestriction

⇒ sudo systemctl daemon-reload

⇒ sudo systemctl restart kube-apiserver

⇒ sudo systemctl status kube-apiserver

/etc/kubernetes/manifests/kube-apiserver.yaml

137

apiVersion: v1

kind: Pod

metadata:

creationTimestamp: null

name: kube-apiserver

namespace: kube-system

spec:

containers:

- Command:

- kube-apiserver

- --authorization-mode = Node, RBAC

- --advertise-address = 172.17.0.107

- --allow-privileged = true

- --enable-bootstrap-token-auth = true

- --enable-admission-plugins = NodeRestriction,

image: k8s.gcr.io/kube-apiserver-amd64:v1.11.3

name: kube-apiserver

To add an admission controller, ~~enable~~ update the

enable admission plugins flag on the kube API server service

to add the new admission controller

kube-apiserver.service

--enable-admission-plugins = NodeRestriction, NamespaceProvision

from KubeADM setup

Update the flag within the kube API server manifest file

as shown below.

/etc/kubernetes/manifests/kube-apiserver.yaml

--enable-admission-plugins = NodeRestriction, NamespaceAutoProvision

To disable admission controller plugins

use the "disable admission plugins flag"

Kube-apiserver Service

ExecStart = /usr/local/bin/kube-apiserver

--advertise-address = \$INTERNAL-IP

--allow-privileged = true

--apiserver-count = 3

--authorization-mode = Node, RBAC

--enable-admission-plugins = NodeRestriction, NamespaceAutoProvision

--disable-admission-plugins = DefaultStorageClass

image: k8s.gcr.io/kube-apiserver-amd64:v1.11.3

name: kube-apiserver

⇒ Once updated, the next time we run the command to provision a pod in a namespace that does not exist yet

⇒ The request goes through authentication, then authorization & then the Namespace Auto Provision Controller at which point it realizes that the namespace doesn't exist so it creates namespace automatically & the request goes through successfully to create the pod

⇒ If you list the namespaces now, you'll see the blue namespace is automatically created

kubectl run nginx --image nginx --namespace blue
pod/nginx created

| Kubectl get namespaces | | |
|------------------------|--------|-----|
| NAME | STATUS | AGE |
| blue | Active | 3m |
| default | Active | 23m |
| kube-public | active | 24m |
| kube-system | Active | 24m |

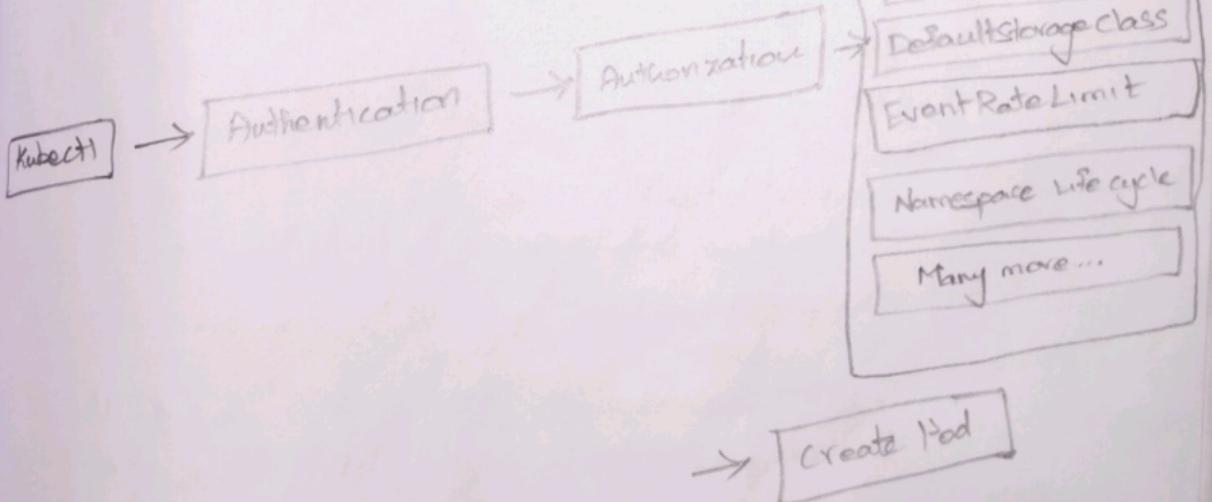
→ So, this is one eg. of or has an "admission controller" works

→ It cannot only validate and rejects request from user, it can also perform operations in backend (or) change the request itself

NOTE:-

NameSpace Auto Provision and NamespaceExists at admission controllers are deprecated & now replaced by the Namespace Lifecycle Admission Controller

→ The "Namespace Lifecycle" admission controller will make sure that requests to a non-existent namespace is rejected & that the default namespace, such as default Kube System & kube public cannot be deleted



Since the kube-apiserver is running as pod you
can check the process to see enable & disabled plugins

```
ps -ef | grep kube-apiserver | grep admission-plugins
```