

rolloutcommand

you can see the status of your rollout by running the command

```
kubectl rollout status deployment/myapp-deployment
```

deployment | myapp-deployment

↓  
Name of the object

To see the history & revision of rollout

```
kubectl rollout history deployment/myapp-deployment
```

This shows you the "revisions" & "history" of our deployment

deployments

"myapp-deployment"

REVISION

CHANGE-CAUSE

1

(none)

2

kubectl apply --filename=deployment-definition  
.yaml --record=true

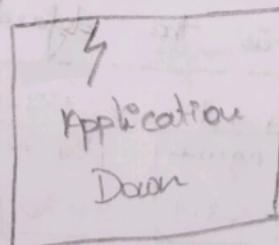
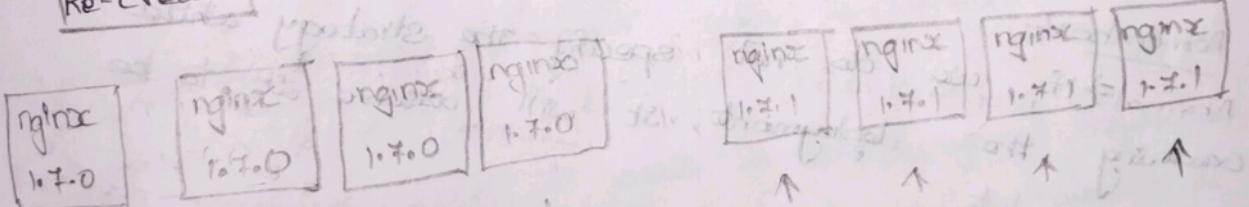
Deployment

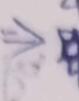
Strategy

⇒ Re-create

⇒ Rolling update

Re-create



⇒  <sup>replica</sup>  
<sup>instance</sup> of web-application is deployed

⇒ One way to upgrade this 4 replicas to newest version is to destroy all of these & create newest versions of application instances

⇒ Meaning:- 1<sup>st</sup> Destroy the 4 running instances & then deploy 4 new instances of the new application version

Problem:- During the period after the older versions are down & before any newer version is up, the application is down and inaccessible to users.

→ This is known as Recreate Strategy, & this is not default deployment strategy

### Rolling update

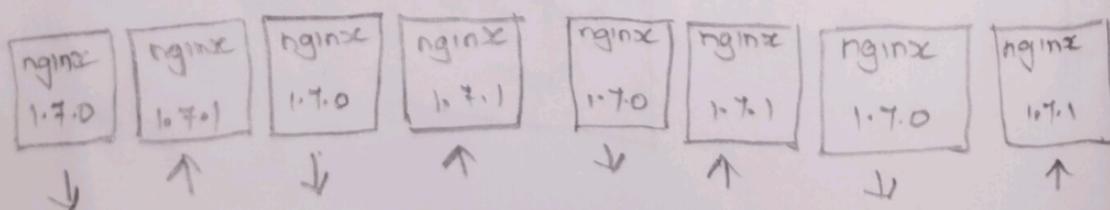
→ In rolling update, where we do not destroy all of them at once, instead we take down the older version & bring up a newer version one by one

→ This way the application never goes down & the upgrade is seamless

#### Remember:-

→ If we do not specify the strategy while creating the deployment, it will assume it to be rolling update

Rolling update is the default deployment strategy



## kubectl apply

165

update - It could be different things such as

updating

\* Application Version [by using the version of docker container used]

\* Updating labels \* updating no. of replicas, etc.

if we have deployment-definition-file, it will be  
easy for us to modify this file by just opening the  
file & edit the image version, save & exit, then kubectl  
apply

> kubectl apply -f deployment-definition.yaml  
deployment "myapp-deployment" configured

once we run the kubectl apply command to apply  
the changes, a new rollout is triggered & a new revision  
of the deployment is created

There is another way to do the same thing  
we can use the "kubectl set image" command to  
update the image of your application

> kubectl set image deployment/myapp-deployment

deployment "myapp-deployment" image is updated  
nginx-container = nginx:1.9.1

Remember:-

Doing it this way will result in the deployment definition  
file having a different configuration, so we may have to be  
cautious when using the same definition file to make changes  
in future

### 3.28 Difference b/w the recreate & Rolling update

⇒ The difference can also be seen, when you view the deployments in detail

#### Kubectl describe deployment myapp-deployment

→ You will notice when the recreate strategy was used the events indicate that the old replicaset was scaled down to zero first & then the new replicaset scaled up to 5.

Name: myapp-deployment

Namespace: default

CreationTimestamp: 2023-07-10T10:45:28Z

Labels: app.kubernetes.io/name=nginx, app.kubernetes.io/version=v1.7.1, k8s-app=nginx, name=nginx

Annotations: deployment.kubernetes.io/revision=1

Selectors: app.kubernetes.io/name=nginx, app.kubernetes.io/version=v1.7.1, k8s-app=nginx, name=nginx

Replicas: 5

StrategyType: Recreate

MinReadySeconds: 0

Events:

Type	Reason	Age	From	Message
Normal	ScalingReplicaset	1m	deployment-controller	scaled up replicaset myapp-deployment to 5
Normal	ScalingReplicaset	1m	deployment-controller	scaled down replicaset myapp-deployment to 0
Normal	ScalingReplicaset	5s	deployment-controller	scaled up replicaset myapp-deployment to 5

update

you view the

Rolling Upd-

at a time

was used

was

a replicaset

when Rolling Update strategy was used the old replicaset was scaled down one at a time, simultaneously scaling up the new replicaset one at a time.

Strategy Type: Rolling Update

Events:

Type	Reason	Age	From	To	Message
Normal	Scaling Replica Set	1m	deployment-controller		Scaled up replicaset myapp-deployment to 5

scaled up replicaset  
myapp-deployment-1sdy  
to 2

scaled down replicaset  
myapp-deployment-1sdy  
to 4

scaled up replicaset  
myapp-deployment-1sdy

scaled down replicaset  
myapp-deployment-1sdy  
to 3

scaled up replicaset  
myapp-deployment-1sdy  
to 4

scaled down replicaset  
myapp-deployment-1sdy  
to 2

replicaset

Deployment to 5

replicaset

Deployment to 0

replicaset

Deployment to 5

## Upgrades

Let's look at how a deployment performs an upgrade

⇒ When a new deployment is created to deploy 5 replicas  
it first creates a replicaset automatically, which  
then creates the no. of pods required to meet the  
no. of replicas

When you upgrade your application, the k8s deployment object creates a new replicaset & starts deploying the containers there at the same time, taking down the pods in the old replicaset following a "rolling update" strategy

This can be seen when you try to list the replicates

Kubectl get replicates					
NAME	DESIRED	CURRENT	READY	AGE	
myapp-deployment-lots	0	0	0	22m	
myapp-deployment-lots above pod id, its unique id	5	5	5	20m	

Now, we see the old replica set with zero pods  
to new replicaset with 5 pods

Today we will discuss  
the important topics

### 5.3 Rollback

169

Once you upgrade your application, you realize something is wrong with the new version of the build you used to upgrade.

so you would like to rollback your update.

K8s deployments allow you to rollback to a previous revision.

To undo a change:-

`kubectl rollout undo deployment/myapp-deployment`  
 deployment "myapp-deployment" rolled back  
 ➔ The deployment will then destroy the pods in the new replicaset & bring the older ones up in the old replicaset & your application is back to its older format.

When you compare the output of the kubectl get replicates command before & after the rollback.

Before the rollback:-

➔ The first replicaset had zero pods & new replicaset had 5 pods & this is reversed after the rollback is finished

NAME	kubectl get replicates				AGE
	DESIRED	CURRENT	READY	REPLICAS	
myapp-deployment-1pd>	0	0	0	0	22m
myapp-deployment-1another pd>	5	5	5	5	20m

NAME	kubectl get replicates				AGE
	DESIRED	CURRENT	READY	REPLICAS	
myapp-deployment-1pd>	5	5	5	5	22m
myapp-deployment-2pd>	0	0	0	0	20m

Initially for the first time, if you want to create the deployment object, then use the kubectl create -f deployment-declaration.yaml

**kubectl create -f deployment-declaration.yaml**

To list the created deployment object

**kubectl get deployments**

If you're updating labels, images, no. of replicas etc, in the manifest file, post editing the manifest file, we

**kubectl apply -f deployment-declaration.yaml**

To set the image using imperative command, which is not recommended because the deployment file in local will not get changed, once the object is updated with manifest + imperative command.

**kubectl set image deployment/myapp-deployment**  
nginx=nginx:1.9.1

Status:-

**kubectl rollout status deployment/myapp-deployment**

**kubectl rollout history deployment/myapp-deployment**

Rollback:-

**kubectl rollout undo deployment/myapp-deployment**