

How Kubernetes works

There are 3 types of file.

1) Local file - developed by SDE, stored in Github

2) last applied Configuration file - which will have annotations

3) live object Configuration file.

Local file

→ It is the file created by SDE to design the infrastructure.

→ Stored in Version Control System

→ YAML file.

Last applied Configuration file

→ Initially we create / edit an object using declarative approach (YAML manifest file), a copy of the local file will be created as metadata, in K8s object metadata, which is called last applied Configuration file

→ It is in JSON representation.

→ Once you've created an object using declarative way (e.g. using manifest file, for eg: here it is "pod creation")

Once the pod got created, execute kubectl get pods

Kubectl got pod in <pod-name> -o yaml

Then look for

metadata:

Annotations: `meta.kubernetes.io/last-applied-configuration: {yaml}`

Kubectl.Kubernetes.io/last-applied-configuration: {yaml}

To view the last applied config from the file, we
Kubectl apply show -last-applied -f <deployment.yaml>
we can't create the annotations manually, Kubectl apply will create
by default under metadata
when last-applied-configuration file will be removed?
> If you want to delete annotation manually

Kubectl annotate pod <pod-name> kubectl.k8s.io/last-applied

use "kubectl annotate" to delete annotation configurations-
resources (pods, deployments, services, etc.) which will delete any

7. If you use kubectl replace (or) kubectl create instead of
apply it does not make any change to the annotation whereas
kubectl apply make changes to the annotation.

Live object configuration file - stored in etcd store
> This file has the exact configuration of
the object in cluster
> Additionally, it has "status" of the object

So, if you're making any changes to the k8s object
then k8s will compare all 3 files before applying
the changes

⇒ If you've created an object using declarative way
then "last applied Configuration" file will be created in
k8s object metadata, then post "kubectl apply" command
the object will be created in cluster, then the configuration
for the object will be stored in k8s memory

⇒ If you use Imperative command to create an object
then "last applied configuration file" will not be created

But if you want "last applied Configuration file" for the object created using imperative command, we can do this. The "live-object configuration file" is stored in "etcd"

1) Export config live object to a file

`kubectl get pod nginx -o yaml > nginx.yaml`

2) Remove unwanted fields like status, metadata, uid

3) Apply declaratively

`kubectl apply -f nginx.yaml`

Initially, if you've created objects using declarative way it has the same configuration in all 3 files. And if you make any changes / edit the object using "Imperative command", then the object will get changed in the cluster, and as per the cluster. the "live object configuration file" will get modified automatically, but the changes were not recorded in "last applied configuration file" nor in "local file".

So then in future if you want to make any change to the same object, there will be no records for the changes made through imperative command if you would have also forgotten what brief some of the other team mate have to make change.

For eg:- Initially the object was created with the configuration of (replicas=3) & (nginx=1.18), then later changes made through imperative command for (replicas=5) now the "live object configuration file" has (replicas=5) & (nginx=1.18) in cluster, but in "local" & "last applied configuration file" has ("replicas=3") & ("nginx=1.18")

Now if we make changes as "nginx=1.19" & apply the changes, now the cluster will have "nginx=1.19" & "replicas=3" as the "last applied configuration file" does not have the changes made through Imperative command, so the best practice here is while creating an object deleting an object always because the declarative way of approach, so that all the files will be in sync.

The imperative approach is for one-time deployment & you no longer need this object. So the best practice is to use "Kubectl replace" command.

Kubectl apply vs Kubectl create vs Kubectl replace

Kubectl apply updates the annotation with "last-applied-configuration".

Kubectl create -f file.yaml does not add the annotation.

Kubectl replace -f file.yaml (only applies to the resource without changing the state) overwrites the resource without respecting the annotation.

Kubectl replace -f file.yaml forces the new state by changing the last "if" which means either "last applied" or "last state" of "last configuration file" annotations will be removed immediately after "Kubectl replace" use "Kubectl replace -f file.yaml --force" ⇒ use this to after the object got corrupted.

Points to Remember:-

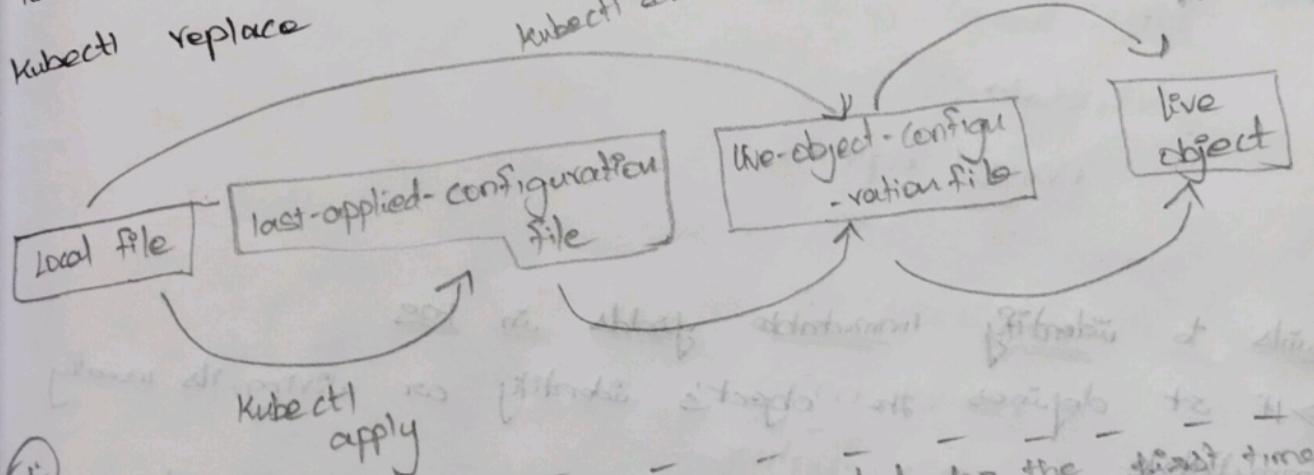
D To delete the annotation:-

Kubectl annotate <resource> <annotation key> -

Kubectl annotate deployment nginx-deployment [example: com/owner] annotation key

2) To prevent future annotations:-
→ When we say "future annotations", we mean avoiding the automatic addition or update of the Kubernetes .90 / last-applied-configuration annotation to a resource.

↳ This annotation is only added/modified by kubectl create, kubectl apply, not by other commands like kubectl create, kubectl replace



↳ `kubectl apply` :- If the object is getting created for the first time "kubectl apply" will create "annotations" in "last-applied configuration file". If you make any changes & execute using "kubectl apply" then annotations will get updated. Hence it is always recommended to use "`kubectl apply`" as it will be helpful for tracking.

`kubectl create` :- This will ignore the "annotations", if a change has been made on local file & executed using "kubectl create", then the annotations will not get updates & if the "object" exists already then it will say "object already exists" & does not make any changes to the annotations in "last-applied-configuration file".

`kubectl replace` :- This will delete/remove the annotation in "last-applied-configuration file" if the object configuration is broken

To view the last-applied-configuration file.

`kubectl get object > object-name > -o yaml`

⇒ The above cmd will show the "last-applied-configuration object" with "annotation" in "YAML" format.

* If you want to make any changes/edits in "last-applied-configuration" file always use `kubectl apply` command as it preserves annotations with changes made.

⇒ Always use `kubectl apply` whenever you are making changes in "last-applied-configuration file".

⇒ `kubectl replace` deletes annotations, use it as last option which is used for quick fixes.

Hints to identify immutable fields in k8s

⇒ If `spec` defines the object's identity or wiring, it's usually immutable.

Field type

Selectors (spec.selector)

Defines which pods the controller manages - changing it breaks ownership.

Service wiring

[clusterIP, nodePort, serviceName]

Controls new identity - changing breaks existing.

PVC access & storage class

[accessModes, storageClassName]

Affects how storage is provisioned
- can't change once bound

Pod Volume setup

[volumeMounts, volumes]

Tied to controller file system - can't swap after start

metadata, name, containerNames

Used as keys as indices
internally - renaming breaks things