

F11 F12 PgUp PgDn

62. Node affinity notes about affinity with?

> The primary purpose of "node affinity" feature is to ensure that pods are hosted on particular nodes.

> we've done this previously using "Node Selector". The challenge in "Node Selector" is that we cannot provide advanced expressions like [OR] or [NOT] as node selectors.

- Large OR Medium
- NOT Small

Node-Selector pod definition file

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
    - name: data-processor
      image: data-processor
  nodeSelector: size=large
  size: Large
```

→ The "Node affinity" feature provides us with advanced capabilities to limit the pod placement on specific nodes.

↳ basically set labels to your nodes, then set the pod's labels to "share affinity" with those nodes. If you want to limit it to a specific node, just add the node's name to the labels.

So the simple "node selector" specification
will now look like this, with "node affinity"
although both does exactly the same thing
but instead does "prioritize" on large nodes.
⇒ Place "to" pool on large nodes.

[yaml] in last and previous section

apiVersion:

kind:

metadata:

name: myapp-pool

spec:

containers:

- name: data-processor

image: data-processor

affinity:

nodeAffinity:

{ requiredDuringSchedulingIgnoredDuringExecution }

nodeSelectorTerms:

- matchExpressions:

- key: size

operator: In

values:

- Large

"Required" means
the node affinity values
in "pod definition file"
should match with node label
of not then pod will not be
scheduled on the node

The **operator: In**

ensures the pod will be placed
on a node, whose label size has any value in the list
of values specified here. In this case, it is just "large"

If you think, your pod should be placed on a
"large or medium node", you could simply add the
value to the list of values like this

Similar to the left side Copy the same & in the "affinity" under "values", add list of values - large, - medium

affinity:

node Affinity:

↳ required During Scheduling Ignored During Execution.
node Selector Terms:

- match Expressions:

- Key: size

operator: In

Values:

- Large

- Medium

you could use **not in** operator to say something like

[size not in Small], where "node affinity" will match the nodes with a size not set to small.

node Selector Terms:

- match Expressions:

- Key: size

operator: NotIn

Values:

- small

we know that we have only set the label size to large and medium nodes. The smaller nodes don't even have the label set, so we don't really have to even check the value of the label. As long as we're sure we don't set a label size to the smaller node, we the exists operator will give us the same result. The exists operator will check if the label size exists on the nodes, & you don't need the value section

for that as it does not compare the values.

[There are other operators as well for which we've to check the documentation]

nodeSelectorTerms:

- matchExpressions:

- Key: size

operator: Exists

⇒ when the pods are created, these rules are considered and the pods are placed onto the right nodes.

1. But what if "node affinity" could not match a node with a given expression. In this case what if there is no node with the label called size?

2. say we have the labels and the pods are scheduled, what if someone changes the label on the node at a future point in time?

will the pod will still lay on the node??

This is answered by a property under node affinity which is called "node affinity type"

Node Affinity Types:-

Available:- If the affinity rule matches then the pod will be scheduled else no scheduler want

1) Required During Scheduling Ignored During Execution

2) Preferred During Scheduling Ignored During Execution

planned:-

Required During Scheduling Required During Execution

The type of node affinity defines the behavior of the Scheduler with respect to node affinity & the stages in the life cycle of the pod.

These are 2 states in lifecycle of a pod when considering node affinity.

- 1) During Scheduling
- 2) During Execution.

During Scheduling

→ It is the state where a pod does not exist & is created for the first time.

→ We have no doubt that when a pod is first created the affinity rules specified are considered to place the pods on the eight nodes.

what if with matching labels are not available

For eg: we forgot to label the node as large, that is, where the type of node affinity used comes into play.

If you select,

Required During Scheduling Ignored During Execution,

→ the scheduler will mandate that the pod be placed on a node with a given affinity rules, if it cannot find one, the pod will not be scheduled.

→ This is used in the case, where the placement of the pod is crucial, if the matching node doesn't exist the pod will not be scheduled.

⇒ If the pod placement is less important, than running the workload itself, in that case, you could set it to "preferred".

⇒ And in cases where a matching node is not found, the Scheduler will simply ignore "node affinity" rules & place the pod on any available node.

The 2nd part of property is during execution.

"Execution"

⇒ During execution is the state where a pod has been already running & a change is made in the environment that affects node affinity.

Such as a change in the label of a node.

For e.g.: If an administrator removed the label we set earlier called "size = large" from the node.

Now what would happen to the pods that're running on the node? As you can see the 2 types of node affinity, has this value set to ignored, which means, pods will continue to run and any changes in node affinity will not impact them once they are scheduled.

Type	Scheduling	During Execution
Type 1	Required	Ignored
Type 2	Preferred	Ignored

61

Required During Scheduling Required During Execution

- The new type expected in the jobs will only have a difference b/w in the during execution phase.
- A new option called ~~req~~ ^{ed} during execution is introduced which will evict any pod that are running on nodes that do not meet affinity rules.
- In the earlier examples a pod running on the large node will be evicted or terminated if the label ~~large~~ is removed from the node.

barrier tag set is prepared as primitive data * what links tag value ~~configuration~~ to it (say tag set in set ~~attribute~~ ~~attribute~~ no terms/no tag set in ~~selected~~ store no header, extra bytes of 2