

Sun, 07/2025

71

- Resource Requirements and Limits
- \* Each and every node will have "CPU" & "Memory"
  - \* And the pod will also require "certain amount of CPU" & "memory" to place this pod on a node.
  - \* If the node doesn't have enough memory or CPU free space, then the scheduler won't schedule the pods on those nodes & instead places the pod on one where sufficient resources are available.
  - \* If there is no sufficient resources available on the nodes, the scheduler holds back scheduling to pods & you'll see pod in "Pending State".

Kubectl describe pod <pod-name>				
NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	Pending	0	7m

Events: short out as kindly step by message between nodes  
Reason: failed connection to known location  
Failed Scheduling: No nodes are available that match all of the following predicates: Insufficient CPU

Resource Request:-  
We can specify the [CPU & memory] required for a pod when creating one.  
For e.g. It could be one CPU & one gigabyte of memory This is known as "resource request".  
So the minimum amount of CPU & memory requested by the container have to be specified in the "pod-definition file", when the scheduler tries to place the pod on the node, it uses these numbers to identify a node which has sufficient amount of resources available.

## pod-definition.yaml

```

version: "v1beta3"
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp-color
  labels:
    name: simple-webapp-color
spec:
  containers:
    - name: simple-webapp-color
      image: simple-webapp-color
  ports:
    - containerPort: 8080
  resources:
    requests:
      memory: "4Gi"
      CPU: 2

```

when the pod gets placed on the node, the pod gets a guaranteed amount of resources available for it.

## Resource - CPU

→ 1 CPU can also 100m [m = milli] expressed as 1 AWS vCPU  
 → 1 AWS vCPU can go as low as 1m but not lower than that

One count of resources at maximum scale	Equal
CPU	is transitive

\* 1 AWS vCPU

\* 1 GCP core

\* 1 Azure Core

\* 1 Hyperthread

## Resource - Memory

you could specify 256 mebibyte using the Mi suffix

- (a) specify the same value in memory like this 268435456
- (b) specify the same value in memory like this 268M  
 $256\text{ MiB} = 268435456 \quad (\text{or}) \quad 268\text{M}$

⇒ use suffix G for gigabyte

1 Gi (Gigabyte) = 1,000,000,000 bytes

1 M (Megabyte) = 1,000,000 bytes

1 K (kilobyte) = 1,000 bytes

1 Gi (Gibibyte) = 1,073,741,824 bytes

1 Mi (Mebibyte) = 1,048,576 bytes

1 Ki (kibibyte) = 1,024 bytes

Limits  
If a container is running on a node & by default, a container has no limit to the resources it can consume on a node, so say a container that is part of a pod starts with one CPU on a node, it can go up & consume as much resources as it requires that suffocates the native processes on the node or other containers of resources.

However, we can set the limit for the resources usage on these pods

For e.g.  
If we set a limit to one VCPU to the Container, a container will be limited to consume only one VCPU

from that node.

The same goes with memory.  
For eg:- you can set a limit of 512 megabyte on containers.

\* we can specify the limits under limits section.

apiVersion: v1

kind: Pod

metadata:

name: Simple-webapp-color

labels:

name: Simple-webapp-color

spec:

containers:

- name: Simple-webapp-color

image: Simple-webapp-color

ports:

- ContainerPort: 8080

resources:

requests:

memory: "1Gi"

CPU: 1

limits:

memory: "2Gi"

CPU: 2

⇒ When the pod is created, k8s sets "new limits" for the container.

④ The limits & requests are set for each container within a pod, so if there are multiple containers then each container can have its own requests or limit set for its own.

## Exceeds Limit

75

what happens when a pod tries to exceed resources beyond its specified limit?

### Incase of CPU:

- The system throttles the CPU, so that it does not go beyond the specified limit.
- A Container cannot use more CPU resources than its limit.

### Incase of memory:

- This is not the case with memory.
- A container can use more memory resources than its limit. So if a pod tries to consume more memory than its limit constantly, then the pod will be terminated & you'll see that the pod terminated with an OOM [Out of Memory] error in the logs or in the output of the described command when you run it.

### Default Behaviour

→ By default, K8S does not have a "CPU" or "memory request" or "limit set". So, this means that any pod can consume as much resources as required on any node & suffocate other pods or processes that are running on the node. This is very important to note.

### Behaviour - CPU

Let's see how "CPU requests" & "limits" work. Let's say there are 2 pods requesting for "CPU requests" on the cluster. If no limit is specified, both pods will get the same fraction of the available resources. So, this is not ideal.

NO REQUESTS but limits specified

Requests = limits

In this case, K8S automatically sets requests to the same as limits.

Eg:- "Requests" and "limits" are assumed to be 3. In this case, each pod is guaranteed 3 vCPUs & no more than that as limits are also set to the same.

## REQUESTS AND LIMITS are set

77

→ In this case, where "requests" & "limits" are set,  
→ In this case, each pod gets a guaranteed # of  
CPU requests, which is one vCPU & can go upto  
the limits that are defined ~~whichever ever~~  
but not more.

→ This might look to be the most ideal  
scenario. However the issue is that if pod1 needs  
more CPU cycles for some reason & pod2  
isn't really consuming that many CPU cycles, then  
we don't want to limit pod1 of CPU  
→ we would like to allow pod1 to use the  
available CPU cycles as long as pod2 doesn't need  
really need it.

→ If there're sufficient CPU cycles available in  
the system then why not let the pods use them  
so, we don't want to unnecessarily limit resources  
of CPU cycles. So, this is not really ideal scenario

### ④ Setting REQUESTS but NO LIMITS

In this case, because "request" is set each pod  
is guaranteed "one vCPU". However, because limits're  
not set when available, any pod can consume as  
many CPU cycles as available. But at any point  
in time, if "pod2" requires additional CPU cycles  
or whatever it has requested, then it will be guaranteed  
its requested CPU cycle.

So this is the most ideal setup

## Behaviour - Memory

If 2 pods competing for "memory resources" on the

cluster, then memory is step-step allocated

NO REQUESTS and NO LIMITS

⇒ Without resource (or) limit set

\* One pod can consume all memory resources on the node & prevent the 2nd part of required resource so this is not ideal

NO REQUESTS specified but LIMITS

⇒ we have no requests specified but we do have limits specified,

⇒ In this case, k8s automatically sets requests to the same as limits

→ For eg:- Requests & limits are assumed to be 3 gibibytes in this case & each pod is guaranteed 3 gibibytes & no more as limits is also the same

REQUESTS & LIMITS are Set

Here, "Requests" & "limits" are Set

In this case, each pod gets a guaranteed amount of memory which is one gibibyte & can go up to "the limits defined which is 3 gibibytes, but not more

Request limits are set as per requirement

setting "requests" but "no limits"

→ setting requests but no limits

⇒ Because requests are set, each pod is guaranteed one gigabyte. However because limits are not set when available any pod can consume as much memory as available and if pod 2 requests more memory to free up pod one, the only option available is to kill it, unlike CPU, we cannot throttle memory. Once memory is assigned to a pod, the only way to kind of relieve it is to kill the pod & free up all the memory that're used by it.

12:20 By default, we know that k8s does not have "resource requests" or "limits" configured for pods. But then how do we ensure that every pod created has some default set. This is possible with "limit ranges".

Limit Ranges:

⇒ Limit Ranges can help you define default values to be set for containers in pods that're created without a "request" or "limit" specified in "pod definition file".

⇒ This is applicable at the "namespace" level.

⇒ "Limit Ranges" is "object", so you create a "definition".

Explanation of the fields Meaning

Field Description

default If no limit is specified in the pod, this will be applied

defaultRequest If no request is specified, this is used.

max Prevents users from setting limits higher than this

min Prevents users from setting limits lower than this

Type: container Applies to individual containers inside a pod

## limit-range.yaml

```
apiVersion: v1
kind: LimitRange
metadata:
  name: cpu-resource-constraint
spec:
  limits:
    default:
      cpu: 500m
    defaultRequest:
      cpu: 500m
    max:
      cpu: "1"
    min:
      cpu: 100m
  type: Container
```

⇒ This is the same goes for "Memory"

use "memory" instead of "CPU". Separation between both objects as request limit & limit range as restrictions

NOTE: These "limits" are enforced when a pod is created so, why you create or change a limit range it doesn't affect existing pods.

It'll only affect newer pods that are created after the limit range is created or updated

Is there any way to restrict the total amount of resources that can be consumed by applications deployed in k8s cluster

Eg:- If we've to say that all pods together that shouldn't consume more than this much of

cpu or memory, we can create "quota" at "namespace" level

↳ Resource Quota is "namespace object" that can be created to set hard limits for "requests" & "limits"

resource-quota.yaml

apiVersion: v1

kind: ResourceQuota

metadata:

name: my-resource-quota

spec:

hard:

requests.cpu: 4

requests.memory: 4Gi

limits.cpu: 10

limits.memory: 10Gi

This resource quota limits the total requested CPU & memory to 4 gibibyte & it defines a maximum limit of CPU consumed by all the pods together to be 10 & maximum limit of memory consumed by all the pods together to be 10 Gibibyte

# A quick note on editing pods and Deployments

83

## Edit a POD

Remember, you CANNOT edit specifications of an existing POD other than the below.

\* spec.containers[\*].image

\* spec.initContainers[\*].image

\* spec.activeDeadlineSeconds

\* spec.tolerations

For eg:-

You cannot edit the

\* Environment Variables

\* Service accounts

\* resource limits (all of which we'll discuss later) of a

running pod.

But if you really have to, then we have 2 options.

i) Run the `Kubectl edit pod <podname>`

→ This command will open the pod specification in an editor (vi editor). Then edit the required property. Then you try to save it, you will be denied.

→ This is because we're attempting to edit a field on the pod that is not editable.

# \*spec:Forbidden: pod updates may not change fields other than 'spec.containers[\*].image'.

#

#

#

apiVersion: v1

" /tmp/kubectl-edit-cvrg.yaml" 125L, 7010C

A copy of the file with your changes is saved in a temporary location.

You can delete the existing pod by running the command.

```
kubectl delete pod webapp
```

Then create a new pod with your changes using the temporary file.

```
kubectl create -f /tmp/kubectl-edit-crvq.yaml
```

2) The second option is to extract the pod definition with YAML format to a file using the command.

```
kubectl get pod webapp -o yaml > my-new-pod.yaml
```

Then make the changes to the exported file using an editor (vi editor)

```
vi my-new-pod.yaml
```

Then delete the existing pod

```
kubectl delete pod webapp
```

Then create a new pod with the edited file

```
kubectl create -f my-new-pod.yaml
```

## Edit Deployments

Deployment changes  
With Deployments you can easily edit any field / property of the Pod template.

since the pod template is a child of the deployment specification, with every change the deployment will automatically delete & create a new pod with new changes

so if you are asked to edit a property of a Pod part of a deployment you may do that simply by running the command.

Kubectl edit deployment my-deployment