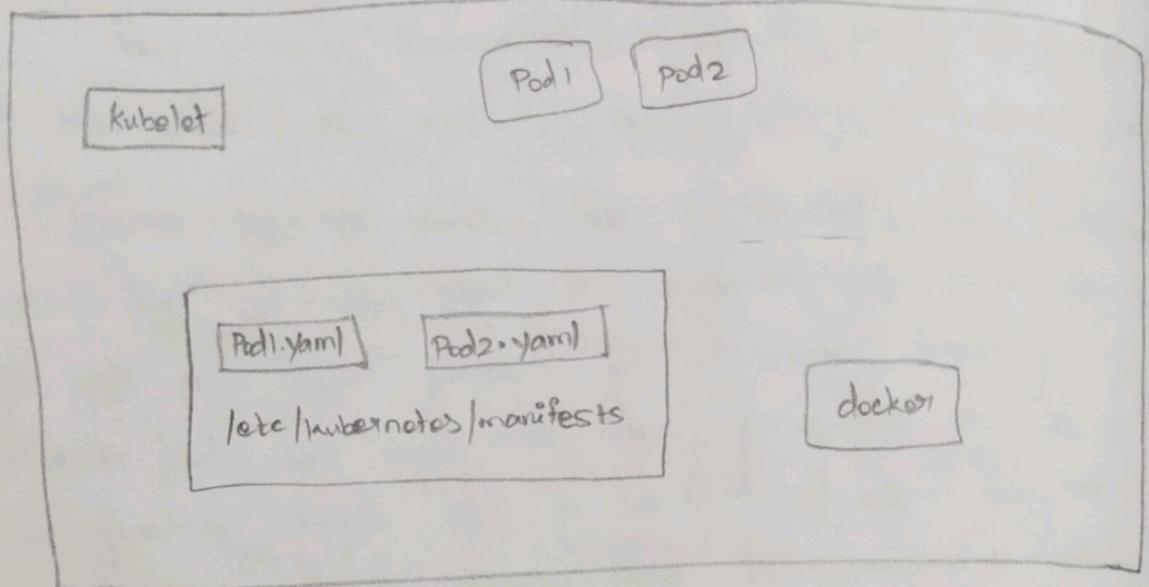


73

Static pods * "control plane" node is usually created using static pods



2:39 what if there is "no master" or "no other node"
if there is "only one node" that has only 'kubelet' installed but which is not part of any cluster

⇒ Is there anything "Kubelet" can do alone, can it operate as independent node, if so, who would provide the instructions required to create these pods

⇒ Kubelet can manage the node independently,
if it has "docker" to run containers. There is no k8s cluster [no master & also other worker nodes]

⇒ The only thing that the kubelet knows to do is create pods, but we don't have an API Server here to provide Pod details, we know that "to create pods" we need the details of the pod in a "pod definition file".

⇒ But how do we provide "pod-definition file" to the "kubelet" without a "kube-apiserver"

⇒ We can configure the "kubelet" to read the "pod definition files" from a directory on the server designated to store information about Pods.

⇒ the kubelet periodically checks this directory for files, reads these files, & creates pods on the host.

host, ^{host} ~~the~~ ^{kubelet} ~~the~~
⇒ Not only it creates the pod, it can ensure that "Pod stays alive", if the application crashes, the "Kubelet" attempts to restart it.

⇒ If you make a change to any of the files within this directory, the kubelet recreates the pod for these changes to take effect.

⇒ If you remove a file from this directory, the pod is deleted automatically.

⇒ These pods created by its own (kubelet) without intervention from the API server or rest of the k8s cluster components are known as "static pods".

2.3 Limitations:-

⇒ we can only create pods this way, you cannot create "Replica sets" or "Deployments" or "Services", by placing a "definition file" in the designated directory. They're all concepts part of the whole k8s architecture that require other cluster plane components like the "replication" & "deployment controllers" etc..

⇒ kubelet only works on pod level & can only understand pods & which is why it is able to create static pods this way.

Static pod

1) pod-manifest-path \rightarrow kubelet.service manifest file

2) config option

So what is that designated folder & how do you
configure it

It could be any directory on the host & the
location of that directory is passed into the
Kubelet as an option while running the Service.

→ The option is named `pod-manifest-path` and here
it is set to `/etc/kubernetes/manifests/folder`.

`kubelet.service`

`ExecStart = /usr/local/bin/kubelet`
or `--container-runtime = remote`
-- `ContainerRuntimeEndpoint = unix:///var/run/containers.sock`
-- `pod-manifest-path = /etc/kubernetes/manifests`

-- `KubeConfig = /var/lib/kubelet/kubeconfig`
-- `network-plugin = cni`
-- `register-node = true`
-- `V=2`

⇒ There is also another way to configure this, instead
of specifying the option directly in the `kubelet.service`
file,

⇒ you could provide a path to another config file
using the `config` option & define the directory path
as `staticPodPath` in that file

⇒ Clusters set up by the "Kube Admin" tool
uses this approach.

Kubelet Service uses string options with values

ExecStart = /usr/local/bin/kubelet // for most
--ContainerRuntime = Remote // for older versions
--ContainerRuntime-endpoint = unix:///var/run/containerd
files in containerd.sock //
--config = kubeconfig.yaml //
--kubeconfig = /var/lib/kubelet/kubeconfig //
--network-plugin = cni //
--register-node = true //
--v = 2 // shows details about state of kubernetes
KubeConfig.yaml is where kubernetes configuration is stored
StaticPodPath: /etc/kubernetes/manifest

If you're inspecting an existing cluster, you should inspect this option of the kubelet to identify the path to the directory. You will then know where to place the definition files for your static pods.

We should know how to "view" & "configure" this option, irrespective of the method used to setup the cluster.

- 1) First, check the option **"pod-manifest-path"** in the "kubelet" service file. If it's not there, then
- 2) Look for the **Config** option to identify the file used as the config file & then within the config file, look for the **"staticPodPath"** option

Either of this should give right path

⇒ Once the static pods are created, you can view them by running the `docker ps` command

⇒ why not "kubectl command", because we don't have the rest of the k8s cluster yet. So the kubectl utility works with the kube-apiserver, since we don't have an API server now, no kubectl utility, which is what we're using `docker` command.

so how does this work, when the node is part of the cluster

when there is an API Server requesting the kubelet to create pods. Can kubelet create both kinds of pods at the same time?

The way the kubelet works is that it can take in requests for creating pods from different inputs

1) First, is through the pod definition files from the static pods folder, as we just saw. They are the ones that start with `pod.yaml`

2) It is through an HTTP API endpoint, & that has the "kube-apiserver" provides input to kubelet only

* The kubelet can provide both kinds of pods

* Static pods & the ones from the API server at the same time

Well in that case, is the API server aware of the static pods created by the kubelet, Yes it is

the "kubelets" are responsible for creating static pods on the node

If we run `kubectl get pods`, as in "if kubectl" who would
get pods on master node, the static
pods will be listed as any other pods

How is this Happening

when the kubelet creates static pod, if it is a
part of a cluster, it also creates a mirror object
in the "kube-apiserver"

so what we see from the "kube-apiserver" is
just a read-only mirror of the pod, you can view
the details of the pod but you cannot edit or
delete it like the usual pods
we can only delete them by modifying the
files from the namespaces manifest place
Note that the name of the pod is automatically
appended with the node name,

In this case **node01**

| NAME | READY | STATUS | RESTARTS | AGE |
|-------------------|-------|-------------------|----------|-----|
| Static-web-node01 | 0/1 | ContainerCreating | 0 | 29s |

So why we have to use static pods
since, static pods are not dependent on the K8S
control plane, you can use static pods to deploy the
control plane components itself as pods on a Node

well, start by uninstalling "kubelet" on all
"master nodes", then create "pod definition files" that uses
docker images of the various control plane components
such as the API Server, Controller, Etcd, etc.

→ place the "manifest file" in the designated manifest folder & the kubelet takes care of deploying the control plane components themselves as pods on the cluster

→ this way, you don't have to download binary config files or worry about the services

Crashing

→ If any of these services were to crash, since it's a static pod, it'll be automatically restarted by the kubelet

→ this is how the kubeadmin tool sets up a k8s cluster which is why when you list the pods in the kube-system ns, you'll see control plane components as pods in a cluster setup by the "kube admin" tool

Kubectl get pods -n kube-system

static pods Daemons DaemonSets

1) Created by the kubelet
2) Deploy control plane component as "Static" Pods

created by Kube-API server
(DaemonSet Controller)

Deploy Monitoring Agents,
Logging Agents on nodes

Ignored by the Kube-Scheduler