

58. Taints and Tolerations

99

Taints are for nodes & Taint \Rightarrow nodes

Tolerations are for pods Tolerations \Rightarrow Pods

Taints & Toleration is a concept where the node has to accept certain pods based on pod's toleration.

\Rightarrow In k8s taints and tolerations work together to control which pods can be scheduled on specific nodes.

\Rightarrow Taints are applied to nodes to indicate that they should not accept certain pods unless those pods have a matching toleration.

\Rightarrow Tolerations, on the other hand, are added to pods to allow them to be scheduled on nodes with specific taints.

Taints:

\Rightarrow Taints are applied to nodes to restrict scheduling. They're essentially "locks" that repel pods that don't have the appropriate "key" (toleration).

\Rightarrow Taints have 3 parts: a key, an optional value, and an effect. Kubectl taint nodes <node> key=value:NoSchedule Block new pods
Effect Evict existing pods
 \Rightarrow The effect determines how the taint affects scheduling.

* NoSchedule:

The K8s Scheduler will not schedule pods that don't have a toleration with for the taint. If the toleration in the pod doesn't match with the taint in the node, then scheduler will not schedule the pod.

* PreferNoSchedule:

The Scheduler will try to avoid scheduling pods without a toleration, but it might do so if there are no other suitable nodes available.

[The system will try to avoid placing the node but it is not guaranteed] if the toleration in the pod doesn't match the taint in the node, then there can be chances that scheduler may or maynot place the pod

No Execute

If a pod without a matching toleration is already running on a node, it will be evicted.

In the toleration, if the pod doesn't match the taint in the pod, then it will not place that pod, if the pod already exists in the nodes then that pod will be evicted.

Tolerations

⇒ Tolerations are applied to pods to allow them to be scheduled on nodes with matching taints.

⇒ they act as the "key" that unlocks the taint on the node.

⇒ Tolerations also have 3 parts:

a key, an optional value, and an effect

⇒ The effect of a toleration must match the effect of the taint at tolerates

In essence, taints and tolerances work like this

1) A node is tainted with a specific taint (key,value, effect)

2) Pods that want to be scheduled on that node must have a toleration that matches the taint's key, value and effect.

3) If a pod doesn't have the correct toleration, it will be repelled from the tainted node by the scheduler.

4) If a pod does have the correct toleration, it can be scheduled on the tainted node.

Tainting the node

* kubectl taint nodes node-name key-value:taint-effect
 if & keep with that **NoSchedule** | prefer **NoSchedule** | No Execute
 it does not affect the scheduling operator
 kubectl taint nodes node1 app=blue:NoSchedule
 at last if the taint in the node & toleration in the pod
 matches then the pod will be scheduled to node1,
 if not then the pod will not get scheduled.

Adding Tolerations to a pod

To add toleration to a pod, first pull the pod-definition file, under the spec section, add a section called **tolerations** make the same values used while creating **the taint**

apiVersion:

kind: Pod

metadata:

name: myapp-pod

spec:

containers:

- name: nginx-container

image: nginx

tolerations:

- key: "app"

operator: "Equal"

value: "blue"

effect: "NoSchedule"

accept = taints of pod
has tolerations

Restrict = node affinity

NOTE :-

Taints & Tolerations are only meant to restrict certain pods nodes from accepting

⇒ Taints and Tolerations doesn't tell the pod to go to a particular node, instead, it tells the node to only accept pods with certain tolerations

→ If your requirement is to restrict a pod to a certain nodes, it is achieved through another concept called **node affinity**

INTERESTING FACT

The scheduler does not schedule any pods on master node, why is that? at node level, when k8s cluster is first setup, a taint is set on the master node automatically that prevents any pods from being scheduled on this node [we can view & change this behavior]

To view:-

```
kubectl describe node kubemaster | grep taint
```

Taints: node-role.kubernetes.io/master: NoSchedule

Add Taint

```
kubectl taint nodes <node> <key>:<effect>
```

Eg:-

```
kubectl taint nodes node1 app=prod
```

Remove Taint:- To untaint a node append - after the effect

```
kubectl taint nodes node1 <key>:<effect> -
```

```
kubectl taint nodes node1 app=prod -
```